

## **DSO553 - Overview**

Overview:

- Big Data Landscape
- What is NoSQL
- MongoDB

References:

- And this [link](#) as well
- [Database Design](#)

### **Instructions:**

Below are a mixture of typical NoSQL questions you may see in an interview. Lets review them together.

---

### **Question 1 - What are the differences between NoSQL and RDBMS?**

#### **RDBMS**

- Primary key: unique and acts as foreign key in other tables for joins
- Predefined schema
- Values are domain-specific, column-homogenous, and not multi-value
- Names of columns within a table are unique
- ACID transactions:  
Example --- I pay someone \$5 for coffee, and \$1000 for rent. Transaction occur separately and one at a time while preserving consistency of the data. Data integrity.
- Vertical Scaling: scaling up, it's about increasing the capacity of your existing hardware

#### **NoSQL**

- No Joins
- Schema-less
- Various data types e.g. Unstructured data(video, images, sensor data)
- Storage Types for given use case e.g. document, graph, column, key-value
- BASE → most important point is that the system is eventually consistent. In addition, data is mostly available and at some date, depending on business logic of implementation, eventually consistent in a distributed system. In fact, there are tools that can predict at which point in time a system can reach consistency.
- Horizontal scaling: Solve a traffic jam by constructing more lanes
- Vertical scaling: Classroom can add more chairs but has a threshold/limit

Web-scale: huge volume of data due to the web

---

## Q2 - What do you understand by "Polyglot Persistence" in NoSQL?

-**3Vs**: Velocity, Volume, Variety (extra: Veracity)

\*More clarity on above:

- Volume: The overall size of the data set
- Velocity: The rate at which the data arrives and also how fast it needs to be processed
- Variety: The wide range of data that the data set may contain—that is, web logs, audio, images, sensor or device data, and unstructured text, among many others types

-**Choosing right tools for the kind of data + goal trying to achieve**

**Polyglot persistence** => Picking the right language for a particular problem can be more productive rather than trying to fit all the aspects of that problem into a single language. Hence, polyglot persistence is the term which is used to define this hybrid approach to data persistence.

### **Pros of Polyglot Persistence**

-Performance. The main advantage is the best performance of the application because each component manages their data using the most appropriate data model and system.

### **Cons:**

- complexity managing data retrieval → requires more resources
- data duplication between data stores/storages/databases

---

## Q3 - List the different kinds of NoSQL data stores.

- document, graph, search, key-value, column
- Unstructured data, eventually consistent, schema-less

### **Key-value:**

Values - are singular in nature for the most part and are not “multi-valued” with exception to Redis

Very performant and fast

Simple queries on keys, not values

### **Column:**

Column-based and great for said simple aggregations e.g. maximums, minimums, sums, averages

No deep analytical queries

No joins

**Graph:**

Nodes, Relationships, Properties

Focused on/built for systems that have lots of relationships (or “joins” in relational system)

Examples: Product recommendations, Social networking

Built for fast traversals -> going from one node to another

**Document:**

-collections ~ tables(sql), documents ~ rows(sql), key:value pairs

-complex data types such as arrays of objects (e.g. list of dictionaries) which is known as the process of denormalization

-Example: product catalogs, metadata, logs

---

**Q4 - What is CAP theorem? How is it applicable to NoSQL systems and data storage ?**

distributed systems: stuff that runs on multiple servers

**CAP Theorem:**

**Consistency:** all nodes/servers in a system have the same view of the data at any time

**Availability:** in the event of some node failure, database remains operational and still sends/receives responses (though responses might be inaccurate or stale)

**Partition tolerance:** database stays operational if system/network fails and where one group of servers are unable to communicate with another group of servers (think about the geographical examples I gave regarding servers not being co-located)

---

**Q5 - What do you mean by eventual consistency in NoSQL stores?**

For banks, consistency is King, so some sort of latency (i.e. lag) penalty is unavoidable and acceptable.

However, for many websites, such as social networks, emails, chatrooms, and certain e-commerce operations (Amazon), this worldwide synchronous consistency is unnecessary. Ultimately, it doesn't matter if my friend in Australia can see my tweet a few seconds before my friend in America. As long as both friends can see the tweet eventually, I'm happy.

Social feed - not everyone sees the same thing at once, but no biggie

Fine-tuning the consistency to meet your needs

EC is OK

-Social media, chatrooms, notifications, emails, weather reporting

EC is bad

-stock market trading, venmo payments, hospital data

---

### **Q6 - What kinds of questions should you be asking when you begin to think about your Big Data implementation?**

-Performance: speed, how fast do you need your system to be? Latency - speed of event occurring

-Availability: do you need 100% uptime?

-Scalability: data storage + how much compute power do you need?

-Flexibility: types of data + changing of data structure AND how quickly can you add more resources

-Cost: what can you afford and what can you justify

---

### **Q7 - What is the Hadoop ecosystem and what necessitated its appearance?**

-distributed, powerful, widely used, can parallelize jobs, data storage mechanism

-moving data from storage to where you can run computations/queries, so that queries run where data exists and not the other way around

-Pre-Hadoop: lots of data loss because raw files were transformed during ETL (extract-transform-load) process

#### **Power of Hadoop:**

Economical

Massive scalability

Reliable

Schema on read

Distributed storage(HDFS) - where all raw files get stored. **Storage**

Distributed processing (MapReduce) - process large data sets using parallel algorithms.

**Processing.**

**Think of Hadoop vs Relational architecture like a restaurant:**

**Relational System** - come to kitchen, give you the cans of meat, tuna, vegetables sliced up, all nice and neat and ready to cook. Easy to do, everything predetermined, but limited

**Hadoop** - Give you the livestock, bag of tomatoes, bag of potatoes, etc. takes much longer and messy. but lots of flexibility and drive to insight

---

**Q8 - What is a Columnar/Wide-Column database and what is a good use case?**

---

**Q9 - What is the biggest case for using Graph database and name examples ?**

---

**Q10 - What is a Key-Value database and what is a good use case?**

---

**Q11 - What is a Document database and what is a good use case?**

---

**Q12 - What are some of the “limits” of Relational Databases**

\*See slides\*

- 1 - Large number of reads & writes
  - 2 - Low latency response times at high volumes
  - 3 - Flexibility in structure and data type
- 

**Q13 - OLTP vs OLAP architecture? What's the difference and can you give a real world example of the differences?**

OLTP - online transaction processing

OLAP - online analytical processing

**OLTP:**

OLTP systems are “classical” systems that process data transactions. They are all around you. In the bank, the ATM or the computer system used by the bank teller to

record a transaction is an OLTP system, usually a database. If you text someone from your smartphone, you are working with another OLTP system. The cash register at your local supermarket runs off another OLTP system, and on it goes.

**OLAP examples:**

systems work with very large amounts of data. Preserving the accuracy and integrity of transactions is not their purpose; this is up to OLTP. OLAP is here to allow us to find trends, crunch numbers, and get the big picture.

**Other Examples:**

OLAP:

- identify sales for each dept each month
- identify top 10 selling books
- find number of classes that have had fewer than 10 students the last 10 years

OLTP:

- update account balance
- enroll in course
- add a book to amazon shopping cart